

CDT Module 1 - R Review

Valerie Bradley

9/19/2018


1. Vectors

- Generate 100 standard normal random variables, and keep only the ones which are greater than 1.
Don't use a loop!

Problem with this is that the number of R.V.s generated


```
all_norm = rnorm(1000)
hist(all_norm)
```

Histogram of all_norm




```
hist(all_norm[all_norm > 1], breaks = 10)
```

Histogram of all_norm[all_norm > 1]



```
truncNorm <- function(n, min){  
  out <- c()  
  
  while(length(out) < n){  
    tmp <- rnorm(n)  
    tmp <- tmp[tmp > min]  
    out <- c(out, tmp) #will copy out each time -- uses more memory. alternately, you could initiali.  
  }  
  
  out[seq_len(n)] #sequence of length n  
}  
  
hist(truncNorm(100,1), breaks = 15)
```

Histogram of truncNorm(100, 1)




- b. Write a function which takes two arguments `n` and `min`, and returns `n` independent random variables from a standard normal distribution truncated below by `min`. Let `min` default to 0.

```
rnrm_trunc = function(n, min = 0){  
  X = rnorm(n)  
  X = unlist(lapply(X, max, min))  
  return(X)  
}
```

```
hist(rnrm_trunc(2000, 1))
```


Histogram of rnorm_trunc(2000, 1)



c. Generate 10k truncated normals with min set at -1 and plot as histogram

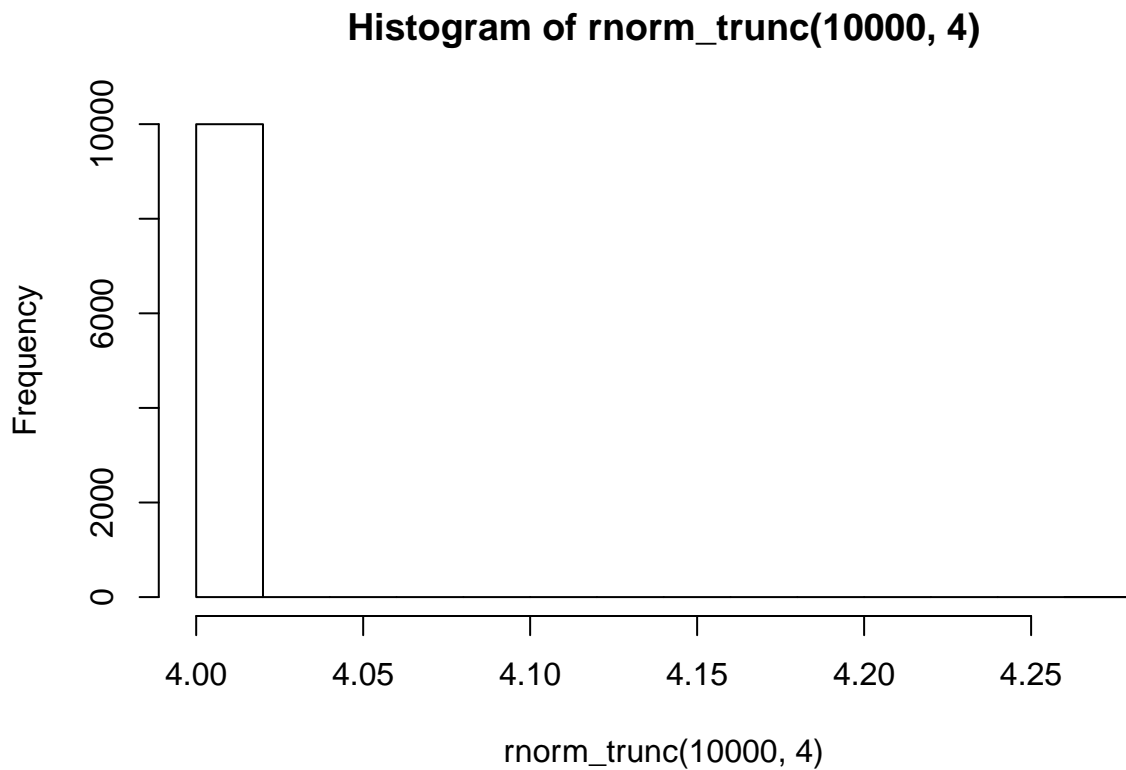
```
samp = rnorm_trunc(10000, -1)
hist(samp, breaks = 10)
```

Histogram of samp



d. what happens if min is large?

```
hist(rnorm_trunc(10000,4))
```



Point here is that this method of rejection sampling is inefficient

2. Data

```
library(MASS)
data(hills)
summary(hills)
```

```
##      dist      climb      time
##  Min.   : 2.000   Min.    : 300   Min.    : 15.95
## 1st Qu.: 4.500   1st Qu.: 725   1st Qu.: 28.00
## Median : 6.000   Median :1000   Median : 39.75
## Mean   : 7.529   Mean    :1815   Mean    : 57.88
## 3rd Qu.: 8.000   3rd Qu.:2200   3rd Qu.: 68.62
## Max.   :28.000   Max.    :7500   Max.    :204.62
```

a. what sort of object is hills?

```
class(hills)
```

```
## [1] "data.frame"
```

b. how many columns?

```
ncol(hills)
```

```
## [1] 3
```

c. Change “Two Breweries” to “Three Breweries”

```

hills[which(rownames(hills) == 'Two Breweries'),]

##           dist climb   time
## Two Breweries   18  5200 170.25

row.names(hills)[which(rownames(hills) == 'Two Breweries')] <- 'Three Breweries'
hills[which(rownames(hills) == 'Two Breweries'),]

## [1] dist   climb time
## <0 rows> (or 0-length row.names)

hills[which(rownames(hills) == 'Three Breweries'),]

##           dist climb   time
## Three Breweries   18  5200 170.25

d. Find the mean time for races iwth a climb greater than 1000ft
with(hills[hills$climb > 1000,], mean(dist))

## [1] 10.41176

```

e. What sort of object is Orthodont? How is it different from hills?

```

library(nlme)
data(Orthodont)
head(Orthodont)

## Grouped Data: distance ~ age | Subject
##   distance age Subject Sex
## 1     26.0   8      M01 Male
## 2     25.0  10      M01 Male
## 3     29.0  12      M01 Male
## 4     31.0  14      M01 Male
## 5     21.5   8      M02 Male
## 6     22.5  10      M02 Male

#class(Orthodont)

```

It's grouped data.

```

head(methods(print))

## [1] "print.abbrev"      "print.acf"          "print.AES"          "print.anova"
## [5] "print.Anova"        "print.anova.lme"

#nlme:::print.groupedData(Orthodont)

```

3. Recursion

Important point here: R is bad at recursion.

The n th Fibonacci number is defined by the recursion $F_n = F_{n-1} + F_{n-2}$ with $F_0 = F_1 = 1$

a. Write a recursive function with argument n which returns the n th Fibonacci number

```

fibonacci = function(n){
  if(n < 0) {
    #print(n)
    F_n = 0
  }else if(n <= 1){ #note this is different than the example given in Recall (they used <=2) - has to

```

```

    #print(n)
    F_n = 1
  }else{
    #print(n)
    F_n = Recall(n-1) + Recall(n-2)
  }
  F_n
}

```

```
fibonacci(2)    #evaluated 3 times
```

```
## [1] 2
```

```
fibonacci(0)    #evaluated 1 time
```

```
## [1] 1
```

```
fibonacci(1)    #evaluated 1 time
```

```
## [1] 1
```

```
fibonacci(3)    #evaluated 5 times
```

```
## [1] 3
```

```

fib_2 = function(n){
  fib_seq = c(1,2)

  if(n < 0) return(0)
  if(n <= 1) return(1)
  if(n == 2) return(2)

  for(i in 3:n){
    fib_seq[i] <- fib_seq[i-1] + fib_seq[i-2]
  }

  return(fib_seq[n])
}

```

```
fib_2(1)
```

```
## [1] 1
```

```
fib_2(2)
```

```
## [1] 2
```

```
fib_2(3)
```

```
## [1] 3
```

```
fib_2(4)
```

```
## [1] 5
```

```
fib_2(5)
```

```
## [1] 8
```

```
fib_2(100)
```

```
## [1] 5.731478e+20
```

4. MCMC

a. $X \sim \text{Gamma}(\alpha, \beta)$ $\alpha, \beta \sim \text{Exp}(1)$

```
#vector of data x
alpha_true = 1
beta_true = 2

x = rgamma(100, shape = alpha_true, rate = beta_true)

# evaluate the posterior distribution of alpha given a vector of data
posterior = function(x, alpha, beta){

  # get prior density
  prior_alpha = dexp(alpha, rate = 1, log = T) # alpha ~ exp(1)
  prior_beta = dexp(beta, rate = 1, log = T) # beta ~ exp(1)

  likelihood = dgamma(x, alpha, beta, log = T)

  posterior = prior_alpha + prior_beta + sum(likelihood)

  return(posterior)
}

posterior(x, alpha = 1.1, beta = 2.3)
```

```
## [1] -7.771917
```

```
posterior(x, alpha = 2, beta = 2)
```

```
## [1] -90.33592
```

b. single metropolis hasting step $\alpha' = \alpha + \sigma Z_1$ $\beta' = \beta + \sigma Z_1$ where Z_1, Z_2 are iid Standard Normal
 $q'(\alpha'|\alpha) N(\alpha, \sigma^2)$

```
step_MH = function(x, alpha, beta, sigma){

  alpha_prime = rnorm(n = 1, alpha, sigma)
  beta_prime = rnorm(n = 1, beta, sigma)

  a = exp(posterior(x, alpha = alpha_prime, beta = beta_prime) - posterior(x, alpha = alpha, beta = beta))
  u = runif(1)
  if(u < a){
    data.frame(alpha = alpha_prime, beta = beta_prime)
  }else{
    data.frame(alpha = alpha, beta = beta)
  }
}

step_MH(x, alpha = 1.1, beta = 2.3, sigma = 0.01)
```

```
##      alpha      beta
## 1 1.09244 2.305277
```

```
run_MH = function(N, x, alpha, beta, sigma){
  chains = data.frame(t = 0, alpha = alpha, beta = beta)
```



```

for(t in 1:N){
  step = step_MH(x, alpha = chains[t,]$alpha, beta = chains[t,]$beta, sigma)
  chains[t+1, ] = c(t, step$alpha, step$beta)
}
return(chains)
}

```


Now, run the full M-H algorithm and plot the resulting chains

```
chains = run_MH(5000, x = x, alpha = 3, beta = 1, sigma = 2/sqrt(length(x)))
```

```

ggplot(chains, aes(x = t, y = alpha)) + geom_line(aes(color = 'alpha')) + geom_line(aes(x = t, y = beta))
ggtitle(paste0("Metropolis-Hastings Sampling for Gamma(alpha, beta)\n", "alpha = ", alpha_true, ", beta = ", beta_true))
geom_hline(yintercept = alpha_true) +
geom_hline(yintercept = beta_true) +
ylab("Parameter estimate")

```



Plot the posterior distributions of the parameters

```


ggplot(chains) + stat_density(aes(x = alpha, fill = 'alpha', alpha = 0.1)) +
stat_density(aes(x = beta, fill = 'beta', alpha = 0.1)) +
geom_vline(xintercept = mean(chains$alpha), linetype = 'dotted') +
geom_vline(xintercept = alpha_true) +
geom_vline(xintercept = mean(chains$beta), linetype = 'dotted') +
geom_vline(xintercept = beta_true) +
ggtitle(paste0("M-H Sampling\nPosterior parameter distributions for Gamma(alpha, beta)\nT = ", nrow(chains)))

```

M-H Sampling

Posterior parameter distributions for Gamma(alpha, beta)

T = 5000



Let's see how this converges over time

```
Ts = seq(250, 5000, by = 250)

all_chains = lapply(Ts, FUN = run_MH, x = x, alpha = 5, beta = 0.3, sigma = 0.01)

all_chains = rbindlist(all_chains)
all_chains = as.data.table(all_chains)
all_chains[, N := rep(Ts, times = Ts + 1)]

all_chains[, .(mean(alpha), .N), N]
```


##	N	V1	N
## 1:	250	4.821162	251
## 2:	500	4.566323	501
## 3:	750	4.143756	751
## 4:	1000	4.000220	1001
## 5:	1250	3.407072	1251
## 6:	1500	3.332075	1501
## 7:	1750	2.999205	1751
## 8:	2000	2.691003	2001
## 9:	2250	2.320869	2251
## 10:	2500	2.398723	2501
## 11:	2750	2.441058	2751
## 12:	3000	2.194616	3001
## 13:	3250	2.181507	3251
## 14:	3500	2.124507	3501

```
## 15: 3750 2.076403 3751
## 16: 4000 2.119446 4001
## 17: 4250 1.959706 4251
## 18: 4500 1.913846 4501
## 19: 4750 1.799522 4751
## 20: 5000 1.607151 5001
```


```
all_chains[, .(mean(beta), .N), N]
```

```
##      N      V1      N
## 1:  250 0.6337948  251
## 2:  500 1.1410063  501
## 3:  750 1.4618715  751
## 4: 1000 1.6336097 1001
## 5: 1250 1.7677558 1251
## 6: 1500 2.1549477 1501
## 7: 1750 2.1744548 1751
## 8: 2000 2.3760333 2001
## 9: 2250 2.4298656 2251
## 10: 2500 2.4278791 2501
## 11: 2750 2.5011539 2751
## 12: 3000 2.6381690 3001
## 13: 3250 2.7352140 3251
## 14: 3500 2.7537821 3501
## 15: 3750 2.8408260 3751
## 16: 4000 3.1044815 4001
## 17: 4250 3.0896072 4251
## 18: 4500 2.7037493 4501
## 19: 4750 2.7185277 4751
## 20: 5000 2.5938450 5001
```

```
ggplot(all_chains[, .(alpha_hat = mean(alpha)), N]) + geom_point(aes(x = N, y = alpha_hat)) + geom_hline
```




```
ggplot(all_chains[, .(beta_hat = mean(beta)), N]) + geom_point(aes(x = N, y = beta_hat)) + geom_hline(y
```



d. Air pollution data.

```
x <- scan("/Users/valeriebradley/Documents/Oxford/Module 1/R review/airpol.txt")
hist(x, breaks = 100, freq = FALSE, main = "Distribution of daily PM2.5 readings in Seattle, 2015", xlab = "PM2.5 readings")
```

Distribution of daily PM2.5 readings in Seattle, 2015




Model data as iid Gamma using $\text{Exp}(1)$ priors from above. Run for 5,000 iterations with starting points $\alpha = 1$, $\beta = 1$.

```
sigmas = c(0.01, 0.02, 0.05)
chains_airpol = lapply(sigmas, FUN = run_MH, N = 5000, x = x, alpha = 1, beta = 1)


chains_airpol = rbindlist(chains_airpol)
chains_airpol = as.data.table(chains_airpol)
chains_airpol[, sigma := sort(rep(sigmas, times = 5000 + 1))]

ggplot(chains_airpol) + geom_line(aes(x = t, y = alpha, group = sigma, color = as.factor(sigma))) +
  ggtitle("Chain convergence of parameter alpha")
```

Chain convergence of parameter alpha




```
ggplot(chains_airpol) + geom_line(aes(x = t, y = beta, group = sigma, color = as.factor(sigma))) +  
  ggtitle("Chain convergence of parameter beta")
```



```
ggplot(chains_airpol, aes(x = alpha, y = beta)) + geom_point() +  
  ggtitle("M-H Exploration of the (alpha, beta) parameter space") +  
  geom_density2d()
```


M-H Exploration of the (α , β) parameter space

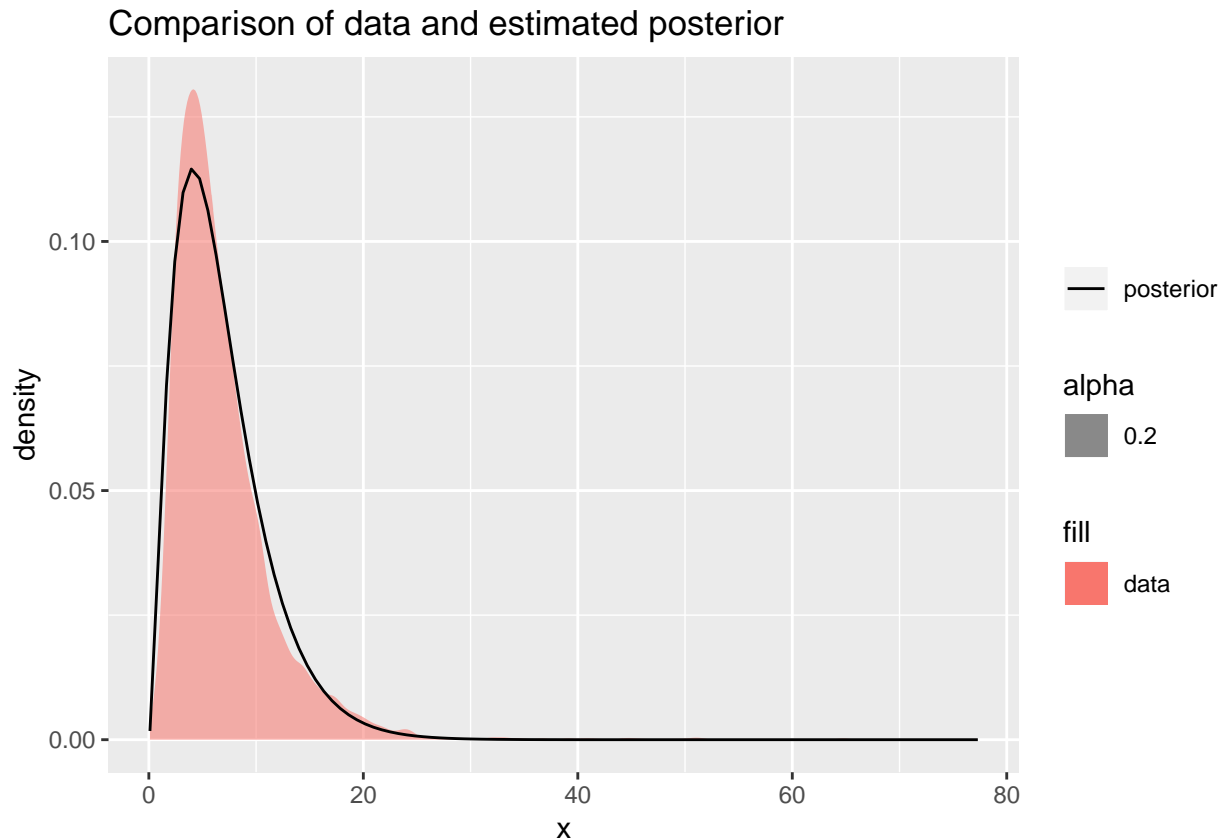


e. Find posterior means of α and β . Plot density of corresponding Gamma distribution over the histogram of the data.

```
chains_airpol[, lapply(.SD, mean), .SDcols = c('alpha', 'beta')]
```

```
##      alpha      beta
## 1: 2.52988 0.3749879
```

```
ggplot(data.frame(x = x)) + stat_density(aes(x = x, fill = 'data', alpha = 0.2)) +
  stat_function(fun = dgamma, args = list(shape = chains_airpol[, mean(alpha)], rate = chains_airpol[,
  scale_colour_manual("", values = 'black') +
  ggtitle('Comparison of data and estimated posterior')
```



5. Methods

```
setClass('biv', representation(x = "numeric", y = "numeric"))
new_df = new('biv', x = rnorm(n = 20), y = rpois(n = 20, lambda = 5))
new_df
```

```
## An object of class "biv"
## Slot "x":
## [1] 2.22093222 0.83426764 -1.11433932 0.27749348 -1.43302676
## [6] -0.79321681 1.06051949 -2.49921643 -1.08725894 -1.22343514
## [11] -1.68157423 -0.28758539 0.45722618 0.04790570 0.33238820
## [16] 0.39156596 0.86217702 0.04231375 1.17919519 0.49887685
##
## Slot "y":
## [1] 9 6 9 4 4 3 7 6 6 1 3 6 4 5 3 4 6 5 2 4
```

Create a print method (that invisibly returns the object)


```
setMethod('print', 'biv', function(x){
  cat(paste0("Bivariate data, ", length(x@x), " entries\n"))
  cat(paste('x : ', paste(x@x[1:min(length(x@x),6)], collapse = " "), "... \n"))
  cat(paste('y : ', paste(x@y[1:min(length(x@y),6)], collapse = " "), "... \n"))
  invisible(x)
})
print(new_df)
```

```
## Bivariate data, 20 entries
```

```
## x : 2.22093222179919 0.834267639367716 -1.11433931599906 0.277493478109035 -1.43302675842844 -0.7932
## y : 9 6 9 4 4 3 ...
```

Construct a plot method for class biv

```
setMethod('plot', 'biv', function(x){
  layout(matrix(c(1,2), 1, 2, byrow = TRUE))
  plot(x = x@x, y = x@y, xlab = 'x', ylab = 'y')
  boxplot(x@x, x@y)
})
plot(new_df)
```



6. Functions

Create a function to return a model matrix from 2 variables

```
getSimpleMatrix = function(x,z){
  if(length(x) != length(z)) stop(print("x and z must have the same number of entries"))

  n_obs = length(x)
  intercept = rep(1, n_obs)
  predictors = c(x,z)
  interactions = x * z

  return(matrix(c(intercept, predictors, interactions), ncol = 4, byrow = F))
}

getSimpleMatrix(x = c(1,2,3,4,5), z = c(6,7,8,9,0))
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    1    6    6
## [2,]    1    2    7   14
## [3,]    1    3    8   24
```

```
## [4,] 1 4 9 36
## [5,] 1 5 0 0
```

Check with model.matrix

```
model.matrix(as.formula(~1+x+z+x:z), data = data.frame(x = c(1,2,3,4,5), z = c(6,7,8,9,0)))
```

```
## (Intercept) x z x:z
## 1 1 1 6 6
## 2 1 2 7 14
## 3 1 3 8 24
## 4 1 4 9 36
## 5 1 5 0 0
## attr(,"assign")
## [1] 0 1 2 3
```

Create a function to return a model matrix from n variables

```
getMatrix = function(...){
  inputs = list(...)

  #check input length
  if(length(unique(unlist(lapply(inputs, length)))) > 1) stop("All inputs must have the same number of elements")

  n_obs = length(x)
  n_vars = length(inputs)

  predictors = matrix(unlist(inputs), ncol = n_vars, byrow = F)
  interactions = apply(combn(x = n_vars, 2), 2, function(x){
    inputs[[x[1]]] * inputs[[x[2]]]
  })

  return(cbind(rep(1, n_obs), predictors, interactions))
}

getMatrix(x = c(1,2,3,4,5), z = c(6,7,8,9,0), y = c(2,4,6,8,10))
```

```
## Warning in cbind(rep(1, n_obs), predictors, interactions): number of rows
## of result is not a multiple of vector length (arg 1)
```

```
## [,1] [,2] [,3] [,4] [,5] [,6] [,7]
## [1,] 1 1 6 2 6 2 12
## [2,] 1 2 7 4 14 8 28
## [3,] 1 3 8 6 24 18 48
## [4,] 1 4 9 8 36 32 72
## [5,] 1 5 0 10 0 50 0
```

7. Mixtures

$X^{(i)} = (X_{i1}, \dots, X_{ik})$ where each X_{ij} is binary A discrete mixture model assumes that each component of the vector $X^{(i)}$ is independent, conditional upon an unknown class label $U_i \in \{1, \dots, l\}$ a. write down the likelihood for one observation $X^{(1)}$, and then for n observations. What are the parameters to be estimated?

$L()$ b.